



**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

**Batch:-** H2-1

**Roll No:-** 16010122151

**Experiment No. 7**

**Title: Classification using support vector machine**

**Aim:** To implement classification using SVM using R libraries

**Expected Outcome of Experiment:**

CO3 : Understand the basic concept and techniques of Machine Learning regression and classification

**Books/ Journals/ Websites referred:**

1. Data Mining Concepts and Techniques Jiawei Han, Michelin Kamber, Jian Pie, 3<sup>rd</sup> edition

---

**Procedure for Implementation in lab :**

1. Select a dataset suitable for classification from UCI data repository or Kaggle.
  - a. Data set used:
  - b. Title:
  - c. Source:
  - d. Number of instances:
  - e. Number of attributes:
  - f. Attribute information :
2. Handle the missing values appropriately
3. Create the SVM model
4. Perform Visualization
5. Calculate prediction accuracy.

**What are SVMs?**

A support vector machine (SVM) is a supervised machine learning algorithm that classifies data by finding an optimal line or hyperplane that maximizes the distance between each class in an N-dimensional space.

**Maximal Margin Classifier**

There are several types of SVM's. The simplest is the\_maximal margin classifier (MMC). Though the MMC is elegant and simple, it cannot be applied to most data sets, since it requires the classes to be separable by a linear boundary.



**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

To visualize an example of separated data, we will generate 40 random data points and assign them to two classes.

```
> set.seed(100)
> # Construct sample data set - completely separated
> x <- matrix(rnorm(20*2), ncol = 2)
> y <- c(rep(-1,10), rep(1,10))
> x[y==1,] <- x[y==1,] + 3/2
> dat <- data.frame(x=x, y=as.factor(y))
> View(dat)
```



**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

	x.1	x.2	y
1	-0.50219235	-0.4380900	-1
2	0.13153117	0.7640606	-1
3	-0.07891709	0.2619613	-1
4	0.88678481	0.7734046	-1
5	0.11697127	-0.8143791	-1
6	0.31863009	-0.4384506	-1
7	-0.58179068	-0.7202216	-1
8	0.71453271	0.2309445	-1
9	-0.82525943	-1.1577295	-1
10	-0.35986213	0.2470760	-1
11	1.58988614	1.4088864	1
12	1.59627446	3.2573756	1
13	1.29836605	1.3620704	1
14	2.23984050	1.3888065	1
15	1.62337950	0.8099857	1
16	1.47068329	1.2782058	1
17	1.11114575	1.6829077	1
18	2.01085626	1.9173233	1
19	0.58618581	2.5654023	1
20	3.81029682	2.4702020	1



**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

	x.1	x.2	y
1	-0.50219235	-0.4380900	-1
2	0.13153117	0.7640606	-1
3	-0.07891709	0.2619613	-1
4	0.88678481	0.7734046	-1
5	0.11697127	-0.8143791	-1
6	0.31863009	-0.4384506	-1
7	-0.58179068	-0.7202216	-1
8	0.71453271	0.2309445	-1

Showing 1 to 8 of 20 entries, 3 total columns

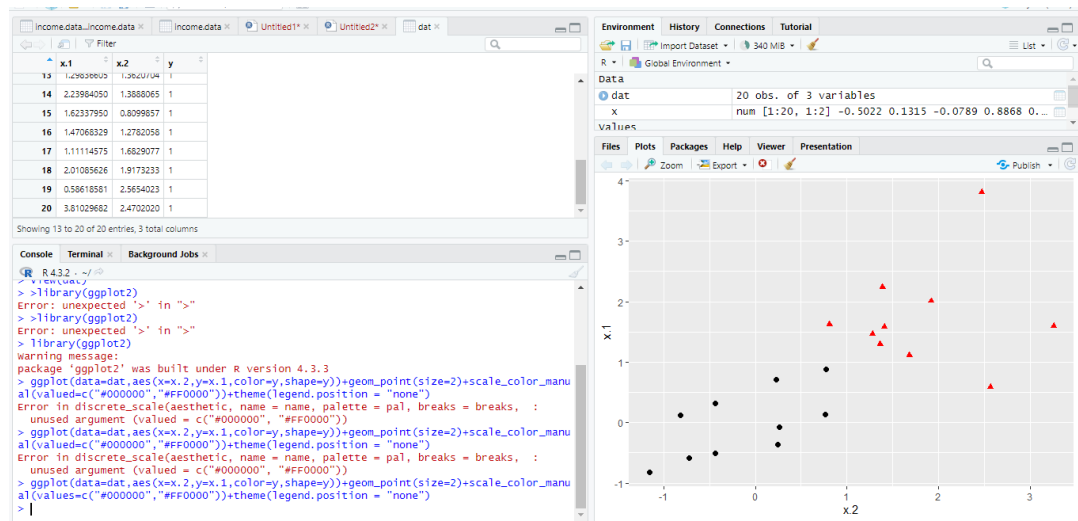
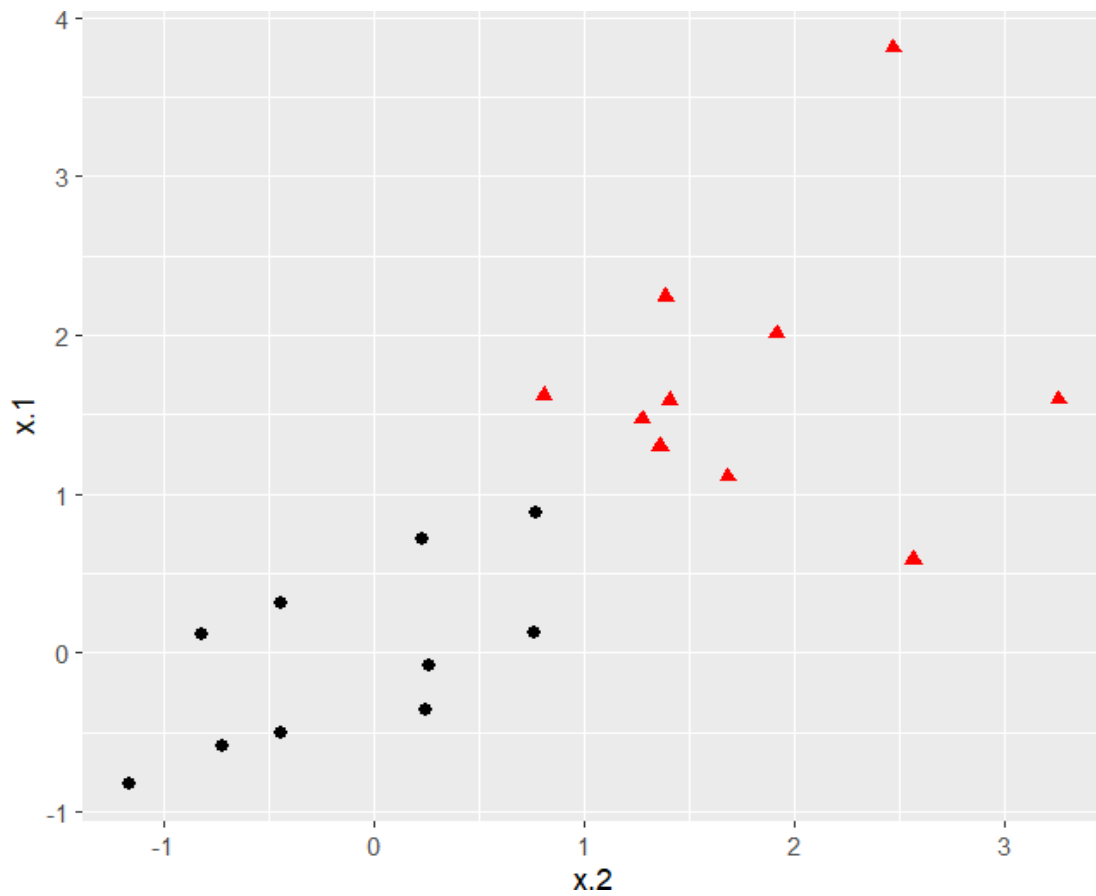
```
R 4.3.2 ~/  
> set.seed(100)  
> x<-matrix(rnorm(20*2),ncol=2)  
> y<-c(rep(-1,10),rep(1,10))  
> x[y==1]<-x[y==1]+3/2  
> dat<-data.frame(x=x,y=as.factor(y))  
> view(data)  
Error in view(data) : could not find function "view"  
> view(dat)  
Error in view(dat) : could not find function "view"  
> View(dat)  
> |
```

**Now we will plot this data**

```
>library(ggplot2)  
> ggplot(data = dat, aes(x = x.2, y = x.1, color = y, shape = y)) +  
+ geom_point(size = 2) +  
+ scale_color_manual(values=c("#000000", "#FF0000")) +  
+ theme(legend.position = "none")
```



**K. J. Somaiya College of Engineering,**  
**Mumbai-77**  
**(A Constituent College of Somaiya Vidyavihar University)**





**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

Upon visual inspection, we can see that infinitely many lines exist that split the two classes. To find the line that maximizes the margin between the classes, we use the `svm()` method from the `e1071` library. We will use 10-fold cross validation.

```
> library(e1071)
> svmfit <- svm(y~., data = dat, kernel = "linear", cross=10)
> svmfit
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cross = 10)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 1
```

Number of Support Vectors: 6

The screenshot shows the RStudio interface. The top pane displays a data frame with columns x.1, x.2, and y. The bottom pane shows the console output, which includes the same R code and results as the text above.

	x.1	x.2	y
13	1.29836605	1.3620704	1
14	2.23984050	1.3888065	1
15	1.62337950	0.8099857	1
16	1.47068329	1.2782058	1
17	1.11114575	1.6829077	1
18	2.01085626	1.9173233	1
19	0.58618581	2.5654023	1
20	3.81029682	2.4702020	1

```
R 4.3.2 ~ /
> svmfit <- svm(y~., data=dat, kernel="linear", cross=10)
Error: unexpected ',' in "svmfit<-svm(y~,"
> svmfit <- svm(y~., data=dat, kernel="linear", cross=10)
Error: unexpected ',' in "svmfit<-svm(y~,"
> svmfit <- svm(y~., data=dat, kernel="linear", cross=10)
> svmfit

Call:
svm(formula = y ~ ., data = dat, kernel = "linear", cross = 10)

Parameters:
SVM-Type: C-classification
SVM-kernel: linear
cost: 1

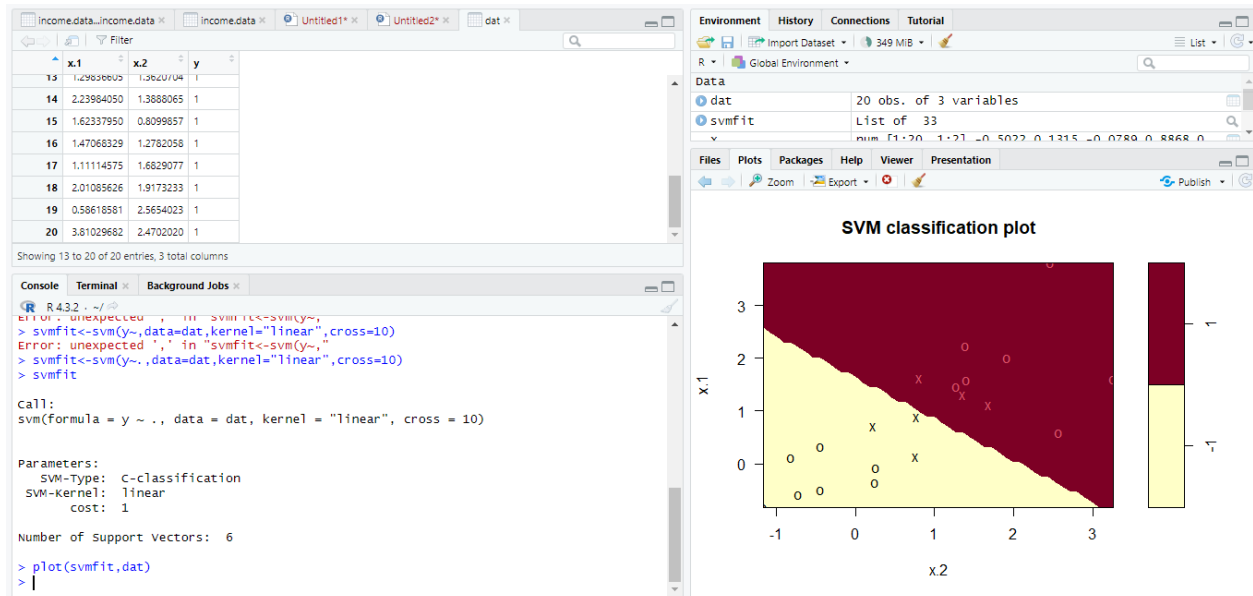
Number of Support Vectors: 6
```



**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

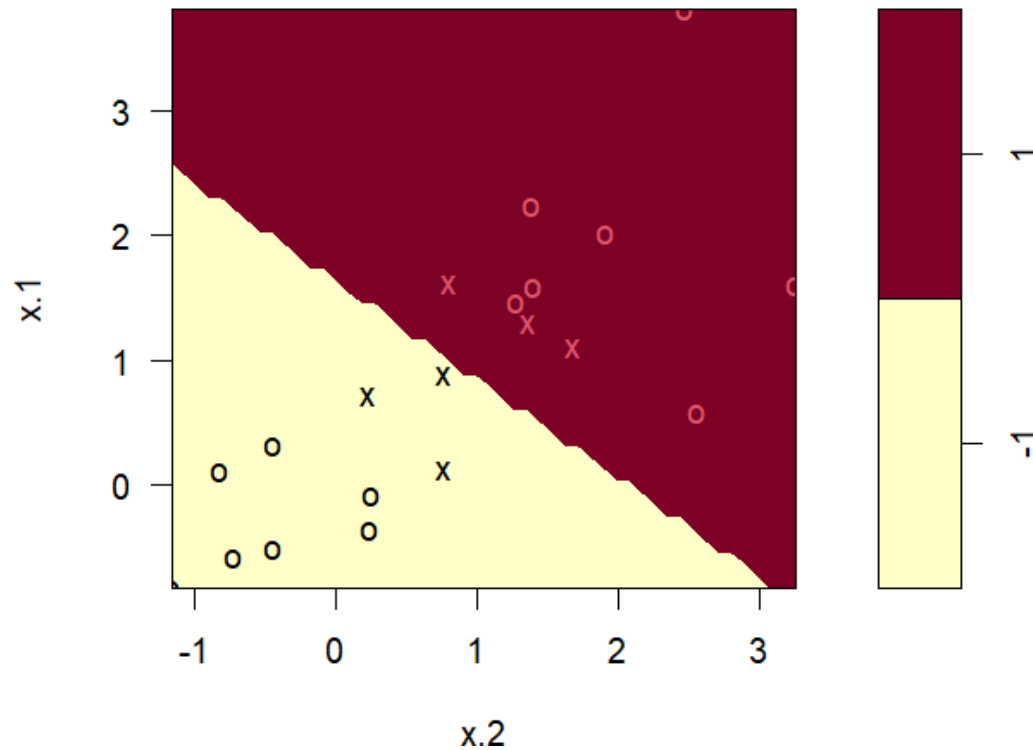
**Visualization:**

**> plot(svmfit,dat)**





### SVM classification plot



Finding the accuracy of the model:

```
> pred <- fitted(svmfit)
> table(pred,y)
```

```
      y
pred -1  1
  -1 10  0
   1  0 10
```

Accuracy = 100% (dataset is too small)





**K. J. Somaiya College of Engineering,**  
**Mumbai-77**  
**(A Constituent College of Somaiya Vidyavihar University)**

income.data...income.data × income.data × Untitled1\* × Untitled2\* × dat ×

Filter

	x.1	x.2	y
13	1.29836605	1.3620704	1
14	2.23984050	1.3888065	1
15	1.62337950	0.8099857	1
16	1.47068329	1.2782058	1
17	1.11114575	1.6829077	1
18	2.01085626	1.9173233	1
19	0.58618581	2.5654023	1
20	3.81029682	2.4702020	1

Showing 13 to 20 of 20 entries, 3 total columns

Console Terminal × Background Jobs ×

R 4.3.2 · ~/

cat:  
svm(formula = y ~ ., data = dat, kernel = "linear", cross = 10)  
  
Parameters:  
  SVM-Type: c-classification  
  SVM-Kernel: linear  
  cost: 1  
  
Number of Support Vectors: 6  
  
> plot(svmfit,dat)  
> pred<-fitted(svmfit)  
> table(pred,y)  
      y  
pred -1  1  
  -1 10  0  
   1  0 10  
> |



**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

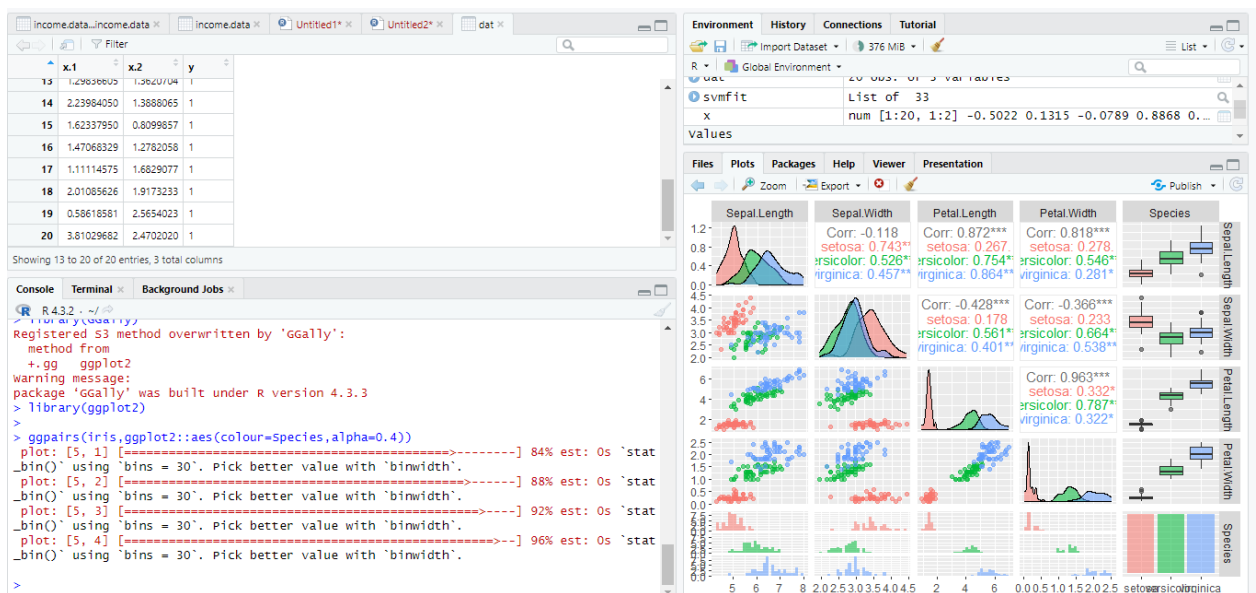
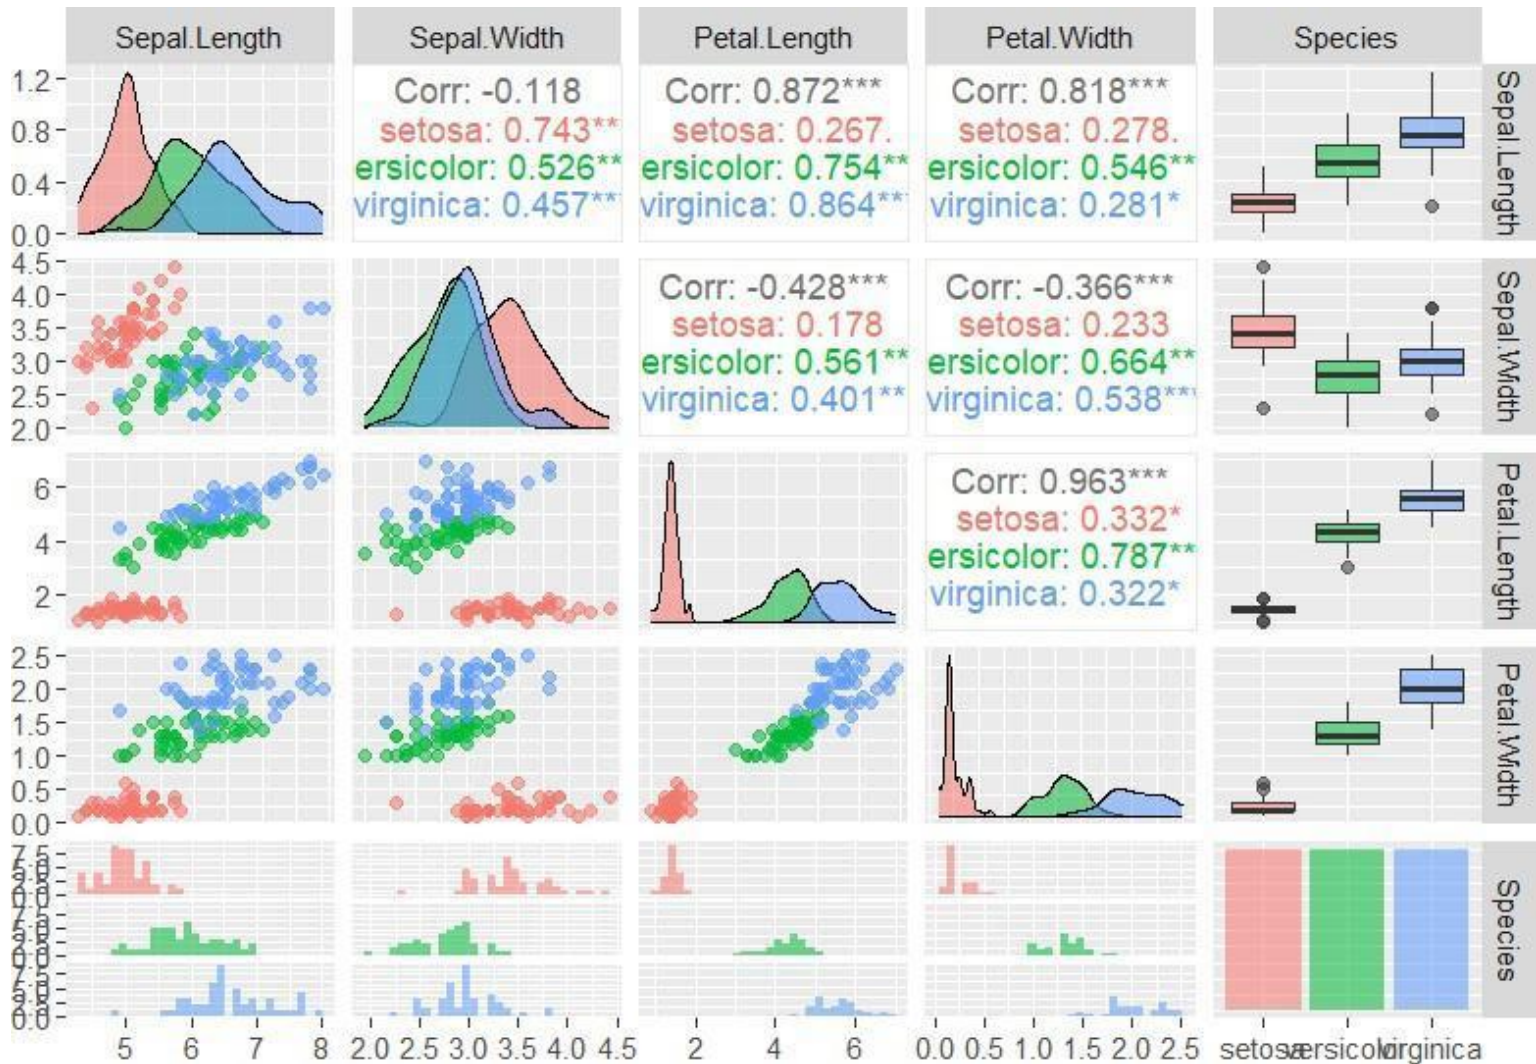
## **Applying SVM to Iris Dataset (4 predictors, 3 classes)**

Since there are 4 predictors, it's not possible to have a simple 2D visualization. Hence we will select only 2 features, just for the sake of visualization. We will use all the four features for training the SVM model.

```
> library(GGally)
Registered S3 method overwritten by 'GGally':
  method from
    +.gg    ggplot2
> library(ggplot2)
> ggpairs(iris, ggplot2::aes(colour = Species, alpha = 0.4))
plot: [5, 1] [=====>-----] 84%
est: 0s `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
plot: [5, 2] [=====>-----] 88%
est: 0s `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
plot: [5, 3] [=====>-----] 92%
est: 0s `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
plot: [5, 4] [=====>-----] 96%
est: 0s `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



**K. J. Somaiya College of Engineering,**  
**Mumbai-77**  
**(A Constituent College of Somaiya Vidyavihar University)**

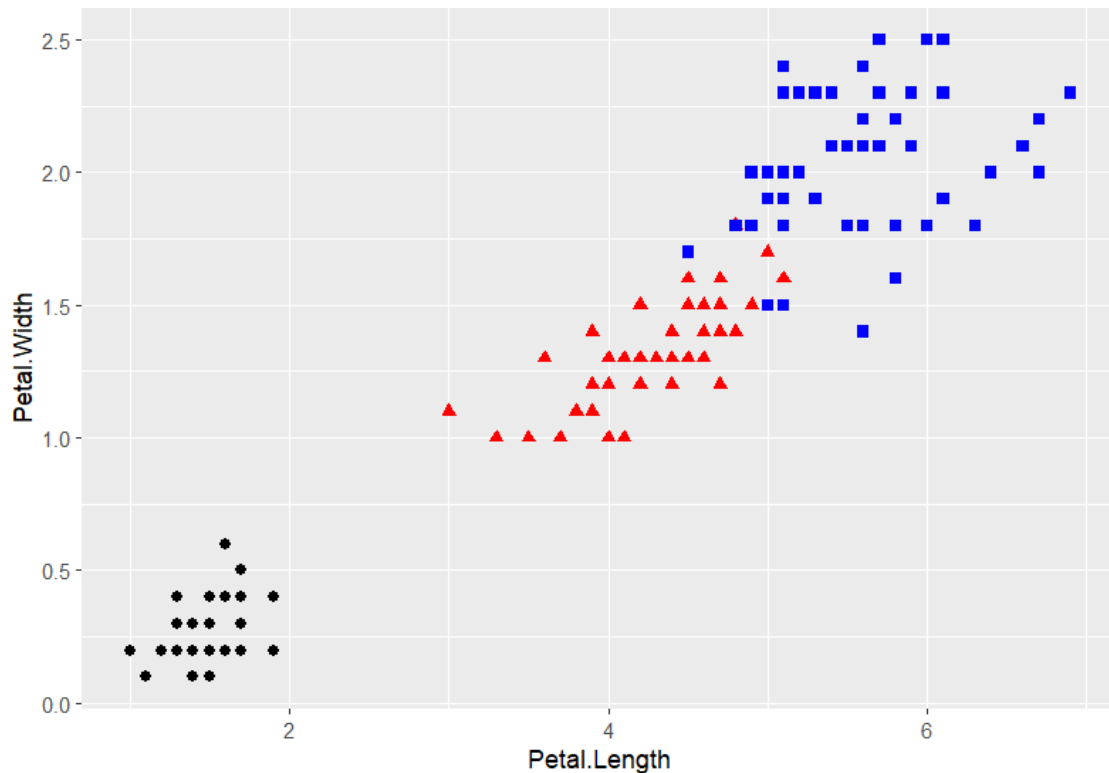




**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

Clearly from the Histograms of **Petal.length** and **Petal.width** that we can clearly separate out Setosa species with very high confidence. However, Versicolor and Virginica Species are overlapped. If we look at the scatterplot of **Petal.Width** vs **Petal.Length**, we can distinctly see a separator that can be drawn between the groups of Species. Looks like we can just use **Petal.Width** and **Petal.Length** as parameters for visualization.

```
> ggplot(data = iris, aes(x = Petal.Length, y = Petal.Width, color = Species, shape = Species)) + geom_point(size = 2)
+ scale_color_manual(values=c("#000000", "#FF0000", "#0000FF"))+ theme(legend.position = "none")
```



```
> lin.svm.iris <- svm(Species~., data = iris, kernel = "linear", cross=10)
> lin.svm.iris
```

Call:

```
svm(formula = Species ~ ., data = iris, kernel = "linear", cross = 10)
```

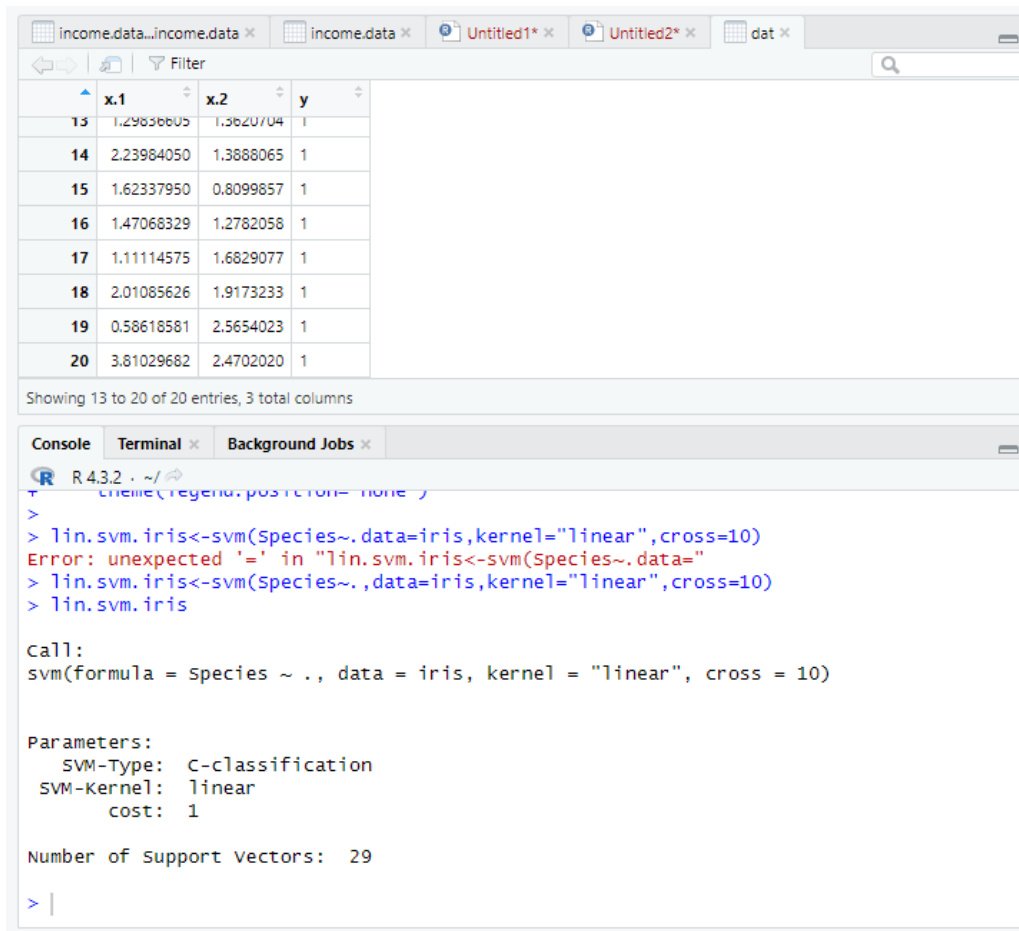
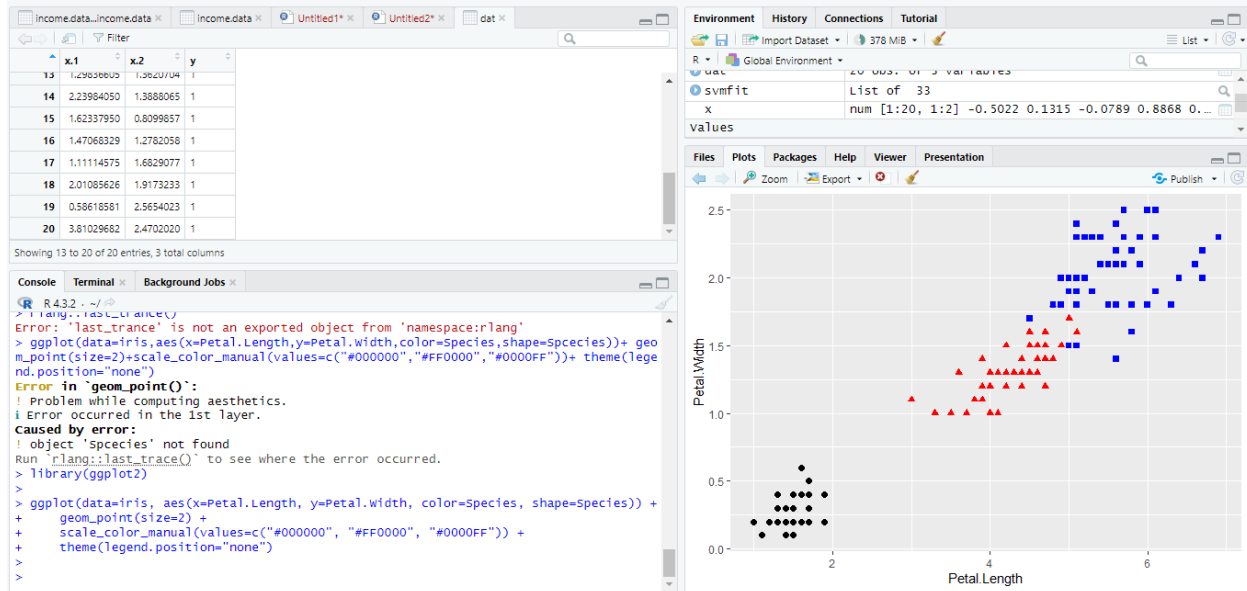
Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 1
```

Number of Support Vectors: 29



**K. J. Somaiya College of Engineering,**  
**Mumbai-77**  
**(A Constituent College of Somaiya Vidyavihar University)**

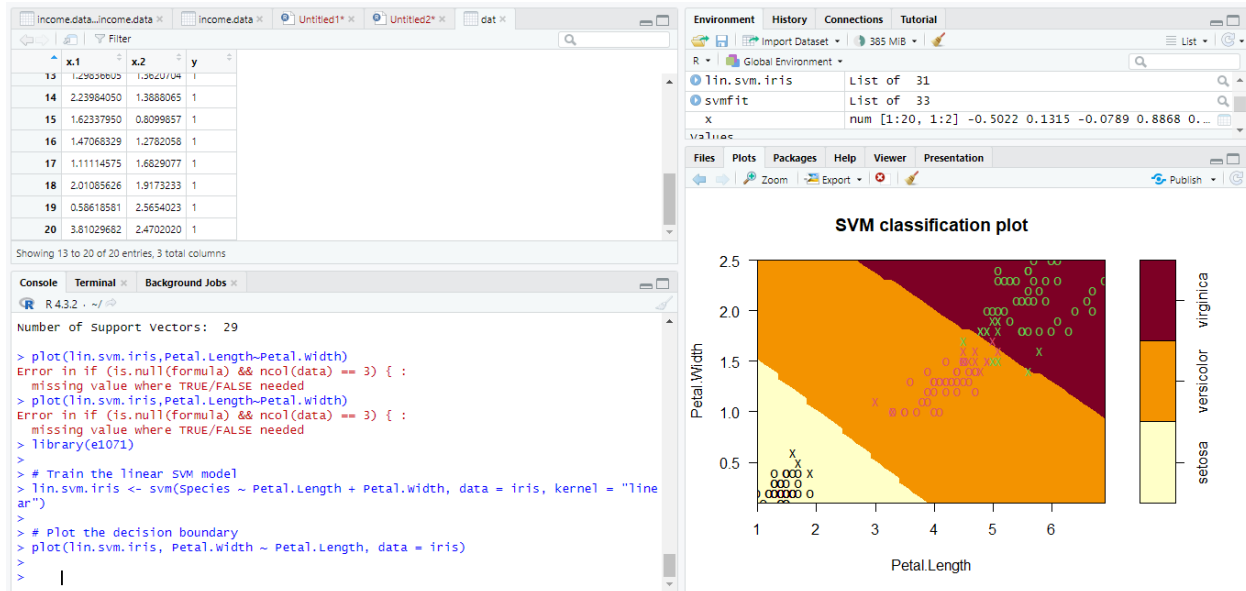




**K. J. Somaiya College of Engineering,  
Mumbai-77**  
(A Constituent College of Somaiya Vidyavihar University)

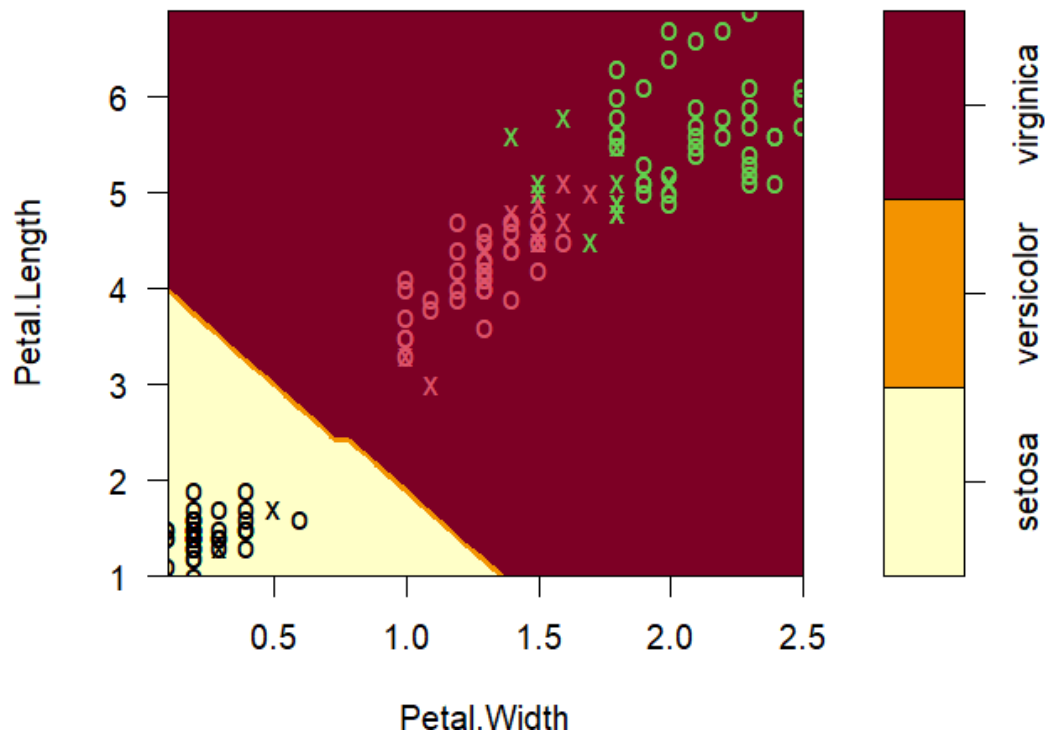
**Visualization:**

```
> plot(lin.svm.iris,iris,Petal.Length~Petal.Width)
```





### SVM classification plot



**Note:** The hyperplane here will be 3-D ( $n$  predictors =  $(n-1)$  D hyperplane)

**Finding the accuracy of the model:**

```
> plot(lin.svm.iris,iris,Petal.Length~Petal.Width)
> lin.svm.iris.pred <- fitted(lin.svm.iris)
> table(lin.svm.iris.pred,iris$Species)
```

lin.svm.iris.pred	setosa	versicolor	virginica
setosa	50	0	0
versicolor	0	46	1
virginica	0	4	49

**Accuracy =  $(50+46+49)/150 = 0.967$**





**K. J. Somaiya College of Engineering,  
Mumbai-77  
(A Constituent College of Somaiya Vidyavihar University)**

The screenshot shows an R Studio window with a data table and a console. The data table has columns x.1, x.2, and y. The console shows R code for training a linear SVM model and plotting the decision boundary.

	x.1	x.2	y
13	1.29836605	1.3620704	1
14	2.23964050	1.3888065	1
15	1.62337950	0.8099857	1
16	1.47068329	1.2782058	1
17	1.11114575	1.6829077	1
18	2.01085626	1.9173233	1
19	0.58618581	2.5654023	1
20	3.81029682	2.4702020	1

Showing 13 to 20 of 20 entries, 3 total columns

```
R 4.3.2 ~/  
installing value where TRUE/FALSE needed  
> library(e1071)  
>  
> # Train the linear SVM model  
> lin.svm.iris <- svm(Species ~ Petal.Length + Petal.width, data = iris, kernel = "linear")  
>  
> # Plot the decision boundary  
> plot(lin.svm.iris, Petal.width ~ Petal.Length, data = iris)  
>  
> plot(lin.svm.iris, iris, Petal.Length~Petal.width)  
> lin.svm.iris.pred<-fitted(lin.svm.iris)  
> table(lin.svm.iris.pred,iris$Species)  
  
lin.svm.iris.pred setosa versicolor virginica  
setosa           50          0           0  
versicolor       0          47          4  
virginica         0          3          46  
>
```

### Conclusion:

**We learnt basic classification using support vector machine (SVM) using R.**

### Post lab Questions:

**Q1. Evaluating the Impact of Different Kernels:**

In this experiment, you used a linear kernel for the SVM. Explore non-linear kernels, such as a polynomial kernel or a radial basis function (RBF) kernel. Train the SVM model again using a non-linear kernel of your choice. Evaluate the performance metrics (accuracy, precision, recall, etc.) for both the linear and non-linear models. Explain the observed differences in performance based on the nature of your data and the chosen kernels.

**Q2. Explain underfitting and overfitting in machine learning. Did you observe any of these when you performed the experiment? How do you handle these issues?**

Underfitting occurs when a model is too simple to capture the underlying patterns in the data, resulting in poor performance on both training and test datasets. Overfitting happens when a model is excessively complex, capturing noise instead of true patterns, leading to excellent performance on the training set but poor generalization to unseen data. Regularization techniques like L1/L2 regularization, cross-validation, and using simpler models can help mitigate these issues.